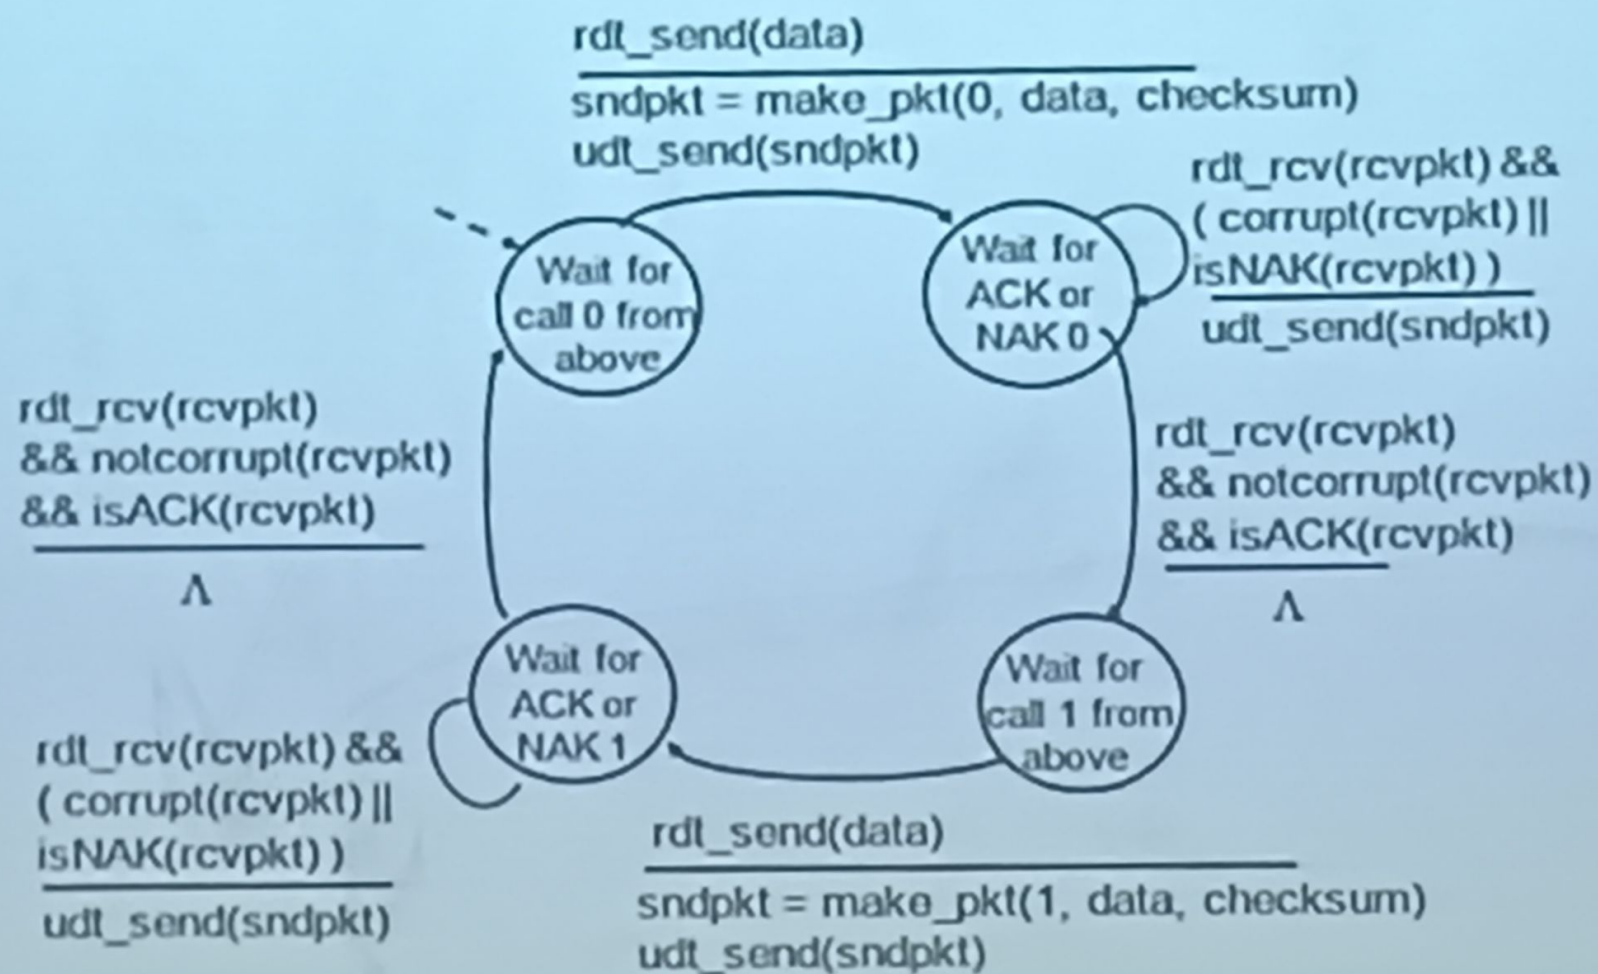


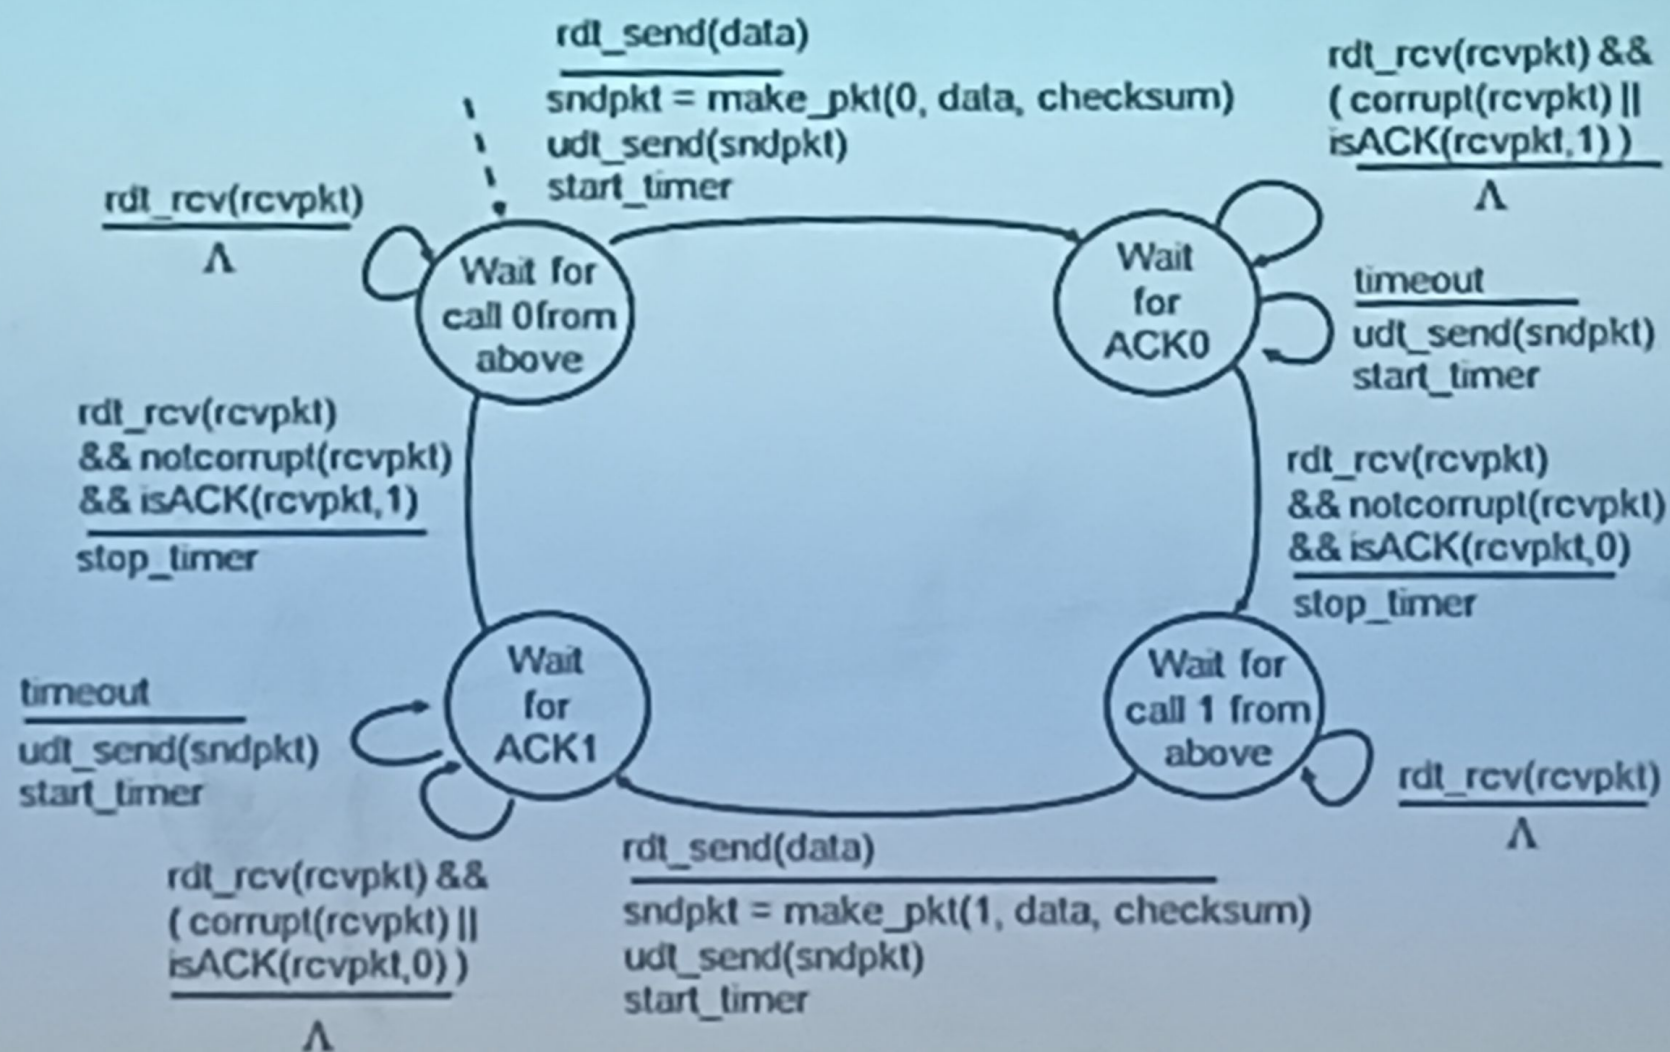
rdt2.1: sender, handles garbled ACK/NAKs



rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK:
retransmit current pkt

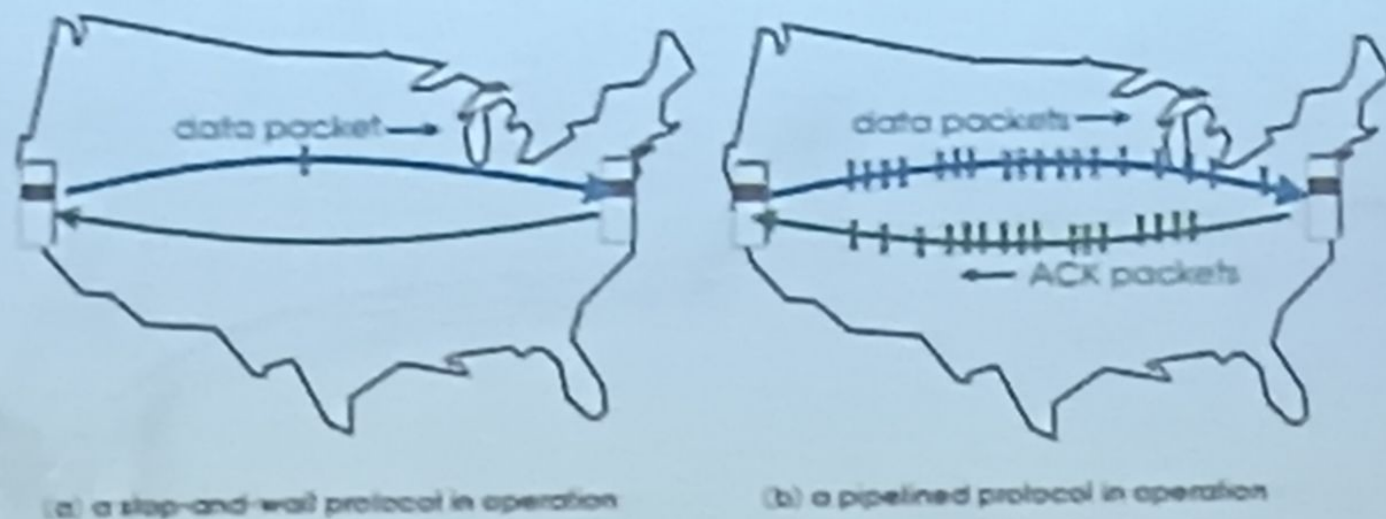
rdt3.0 sender



Pipelined protocols

Pipelining: sender allows multiple, "in-flight", yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver



- Two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

Pipelining Protocols


Go-back-N: big picture:

- Sender can have up to N unacked packets in pipeline
- Rcvr only sends cumulative acks
 - Doesn't ack packet if there's a gap or out of order packet received
- Sender has timer for oldest unacked packet
 - If timer expires, retransmit all unacked packets


Selective Repeat: big pic

- Sender can have up to N unacked packets in pipeline
- Rcvr acks individual packets
- Sender maintains timer for each unacked packet
 - When timer expires, retransmit only unack packet

Send window for Go-Back-N ARQ



a. Send window before sliding



b. Send window after sliding


The send window is an abstract concept defining an imaginary box of size $2^m - 1$ with three variables: S_f , S_n , and S_{size} .

The send window can slide one or more slots when a valid acknowledgment arrives.

Go-Back-N

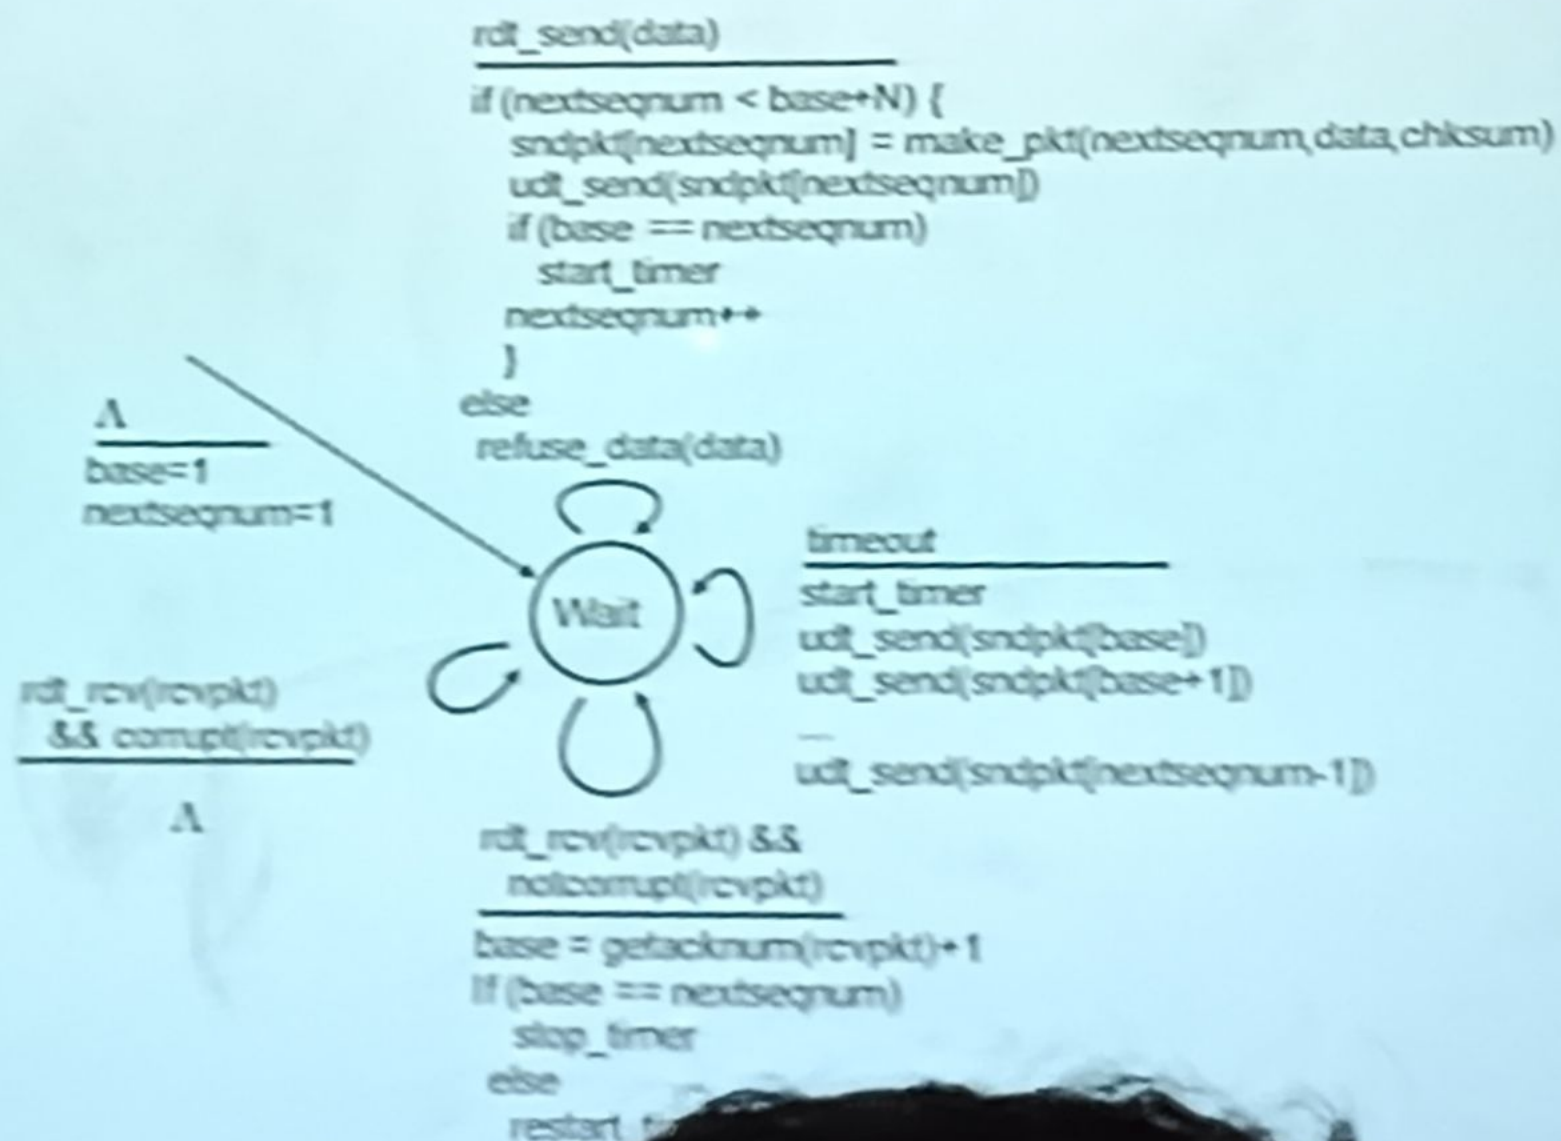
Sender:

- k-bit seq # in pkt header
- "window" of up to N, consecutive unack'd pkts allowed
- send_base: oldest unacked packet
- nextseqnum: smallest unused seqnum

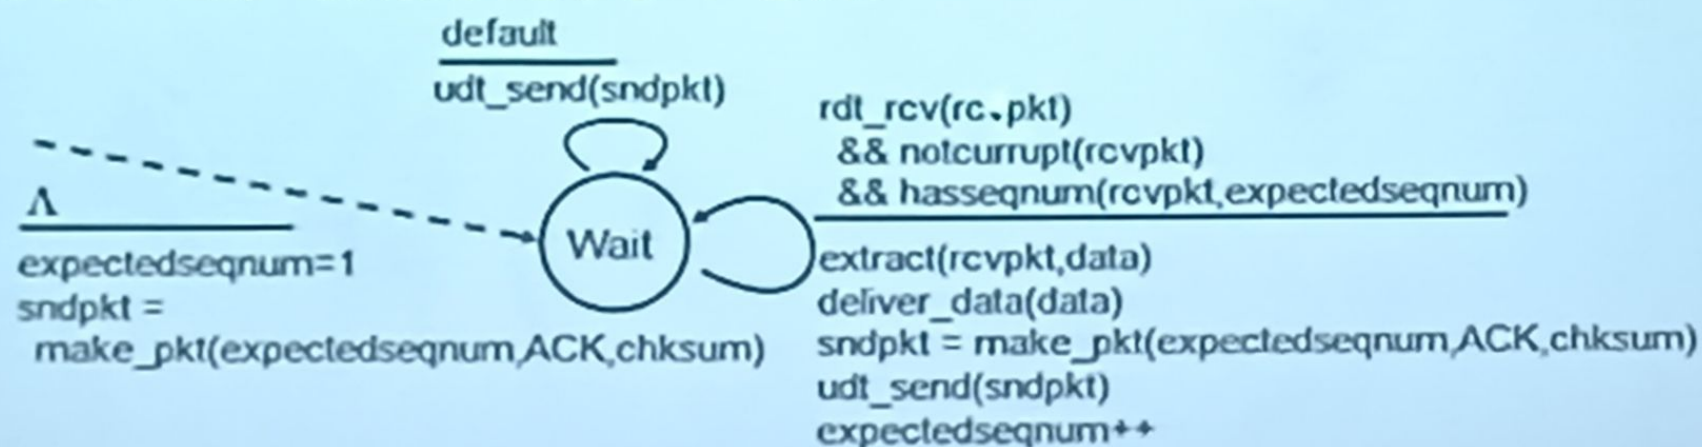


- $\text{nextseqnum} - \text{send_base} \leq \text{Sender Window Size (SWS=N)}$
- $\text{ACK}(n)$: ACKs all pkts up to, including seq # n - "cumulative ACK"
 - may receive duplicate ACKs (see receiver)
- timer for each in-flight pkt
- $\text{timeout}(n)$: retransmit pkt n and all higher seq # pkts in window

GBN: sender extended FSM



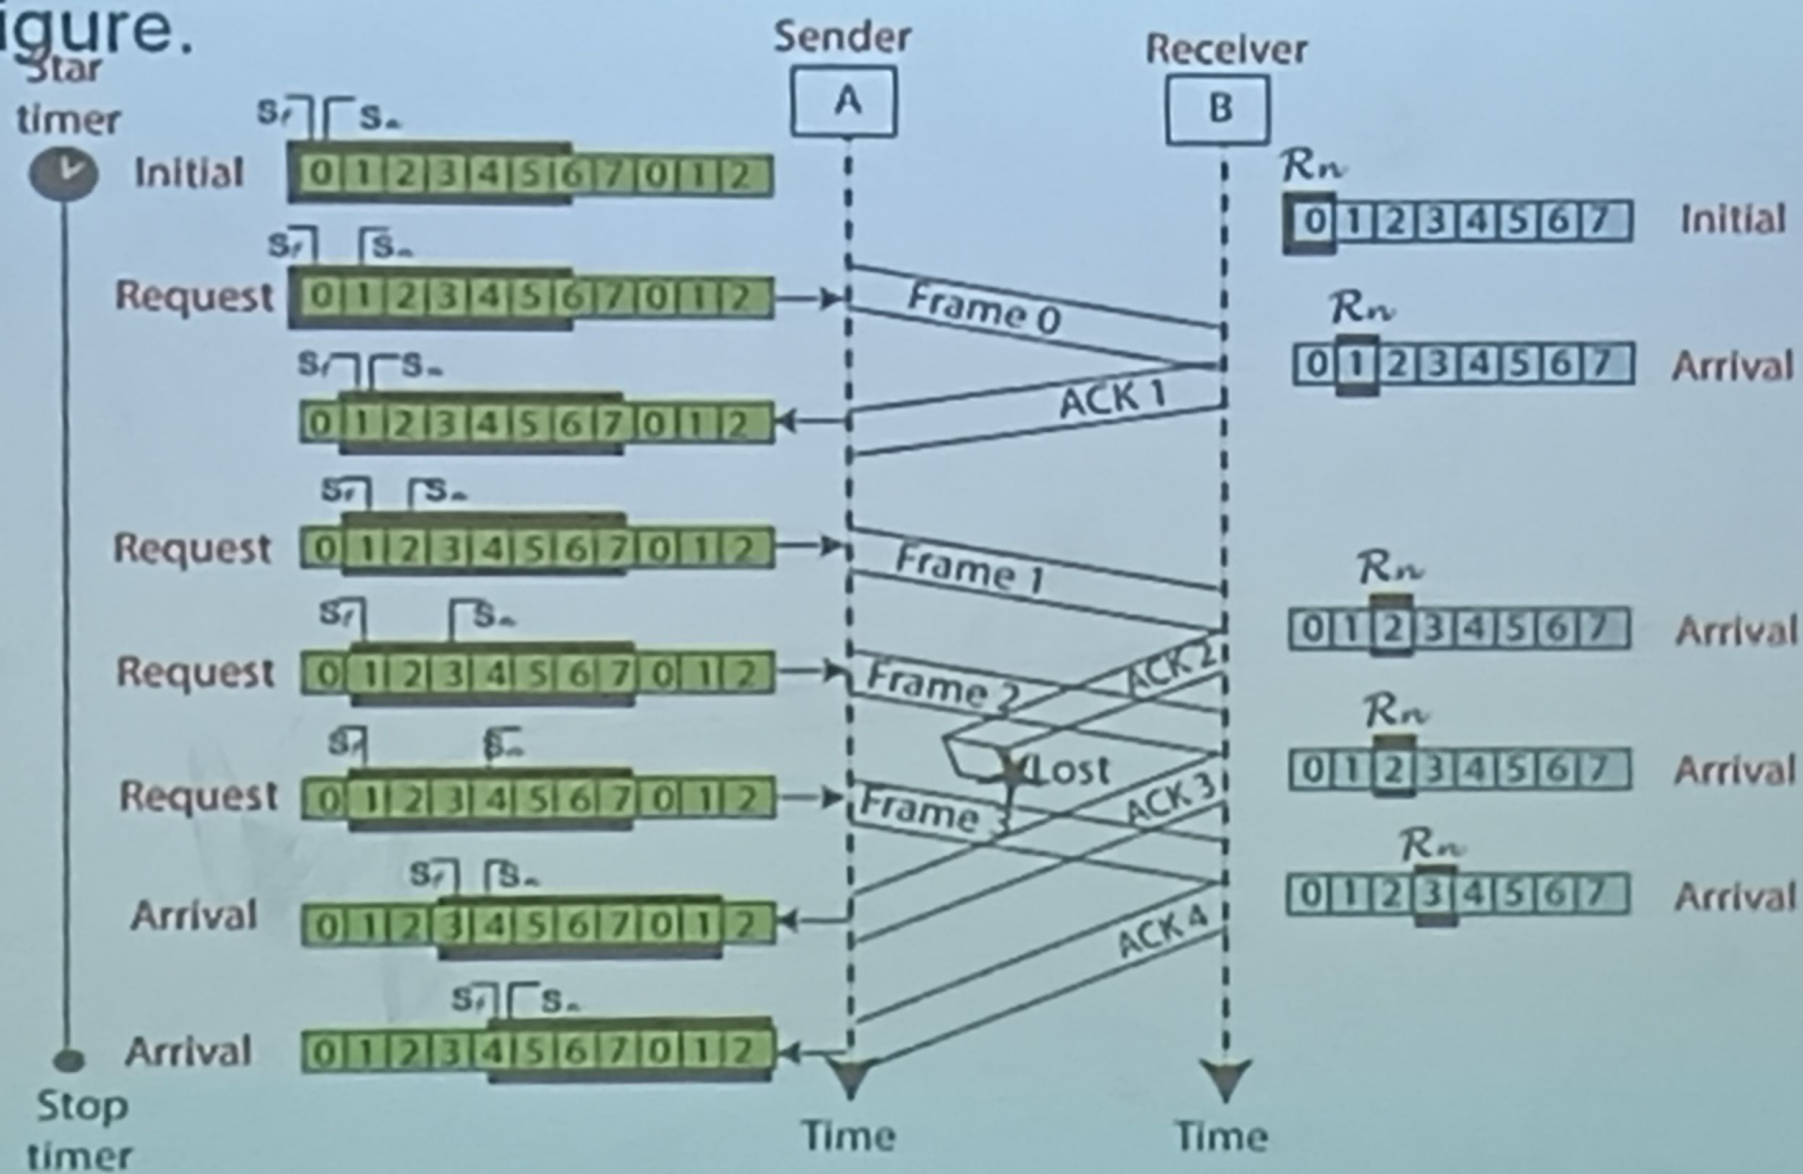
GBN: receiver extended FSM




ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #

- may generate duplicate ACKs
- need only remember **expectedseqnum**
- out-of-order pkt:
 - discard (don't buffer) -> no receiver buffering!
 - Is it bad?
 - Re-ACK pkt with highest in-order seq #

The example of Go-Back-N is shown below in the figure.

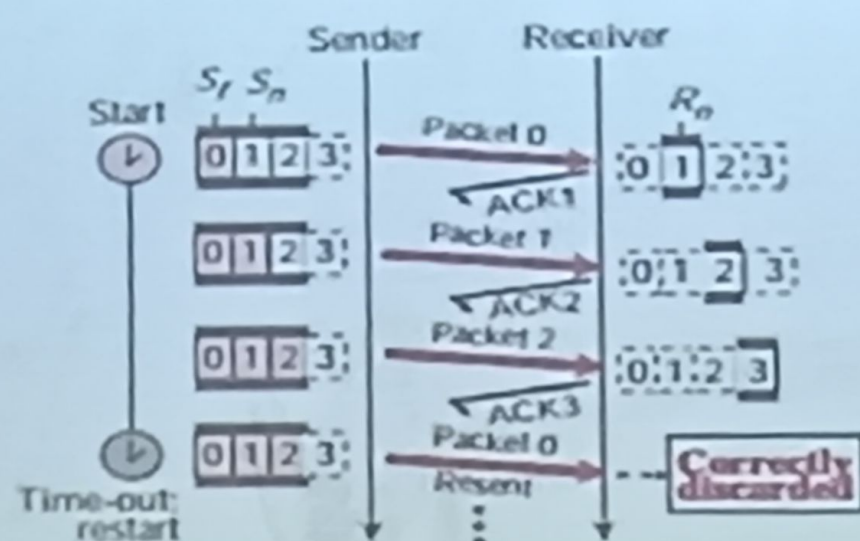


GBN in action

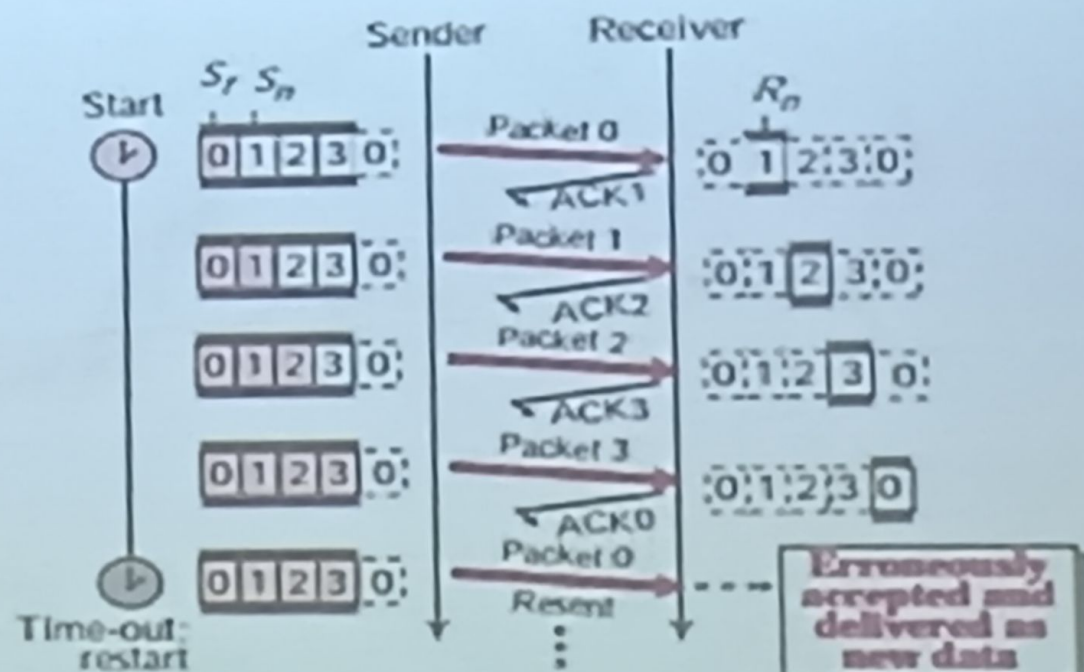


Sender Window size

Send window size for Go-Back-N



a. Send window of size $< 2^m$




b. Send window of size $= 2^m$


Selective repeat: big picture

- Sender can have up to N unacked packets in pipeline
- Rcvr acks individual packets
- Sender maintains timer for each unacked packet
 - When timer expires, retransmit only unack packet

Selective repeat: sender, receiver windows




(a) sender view of sequence numbers




(b) receiver view of sequence numbers

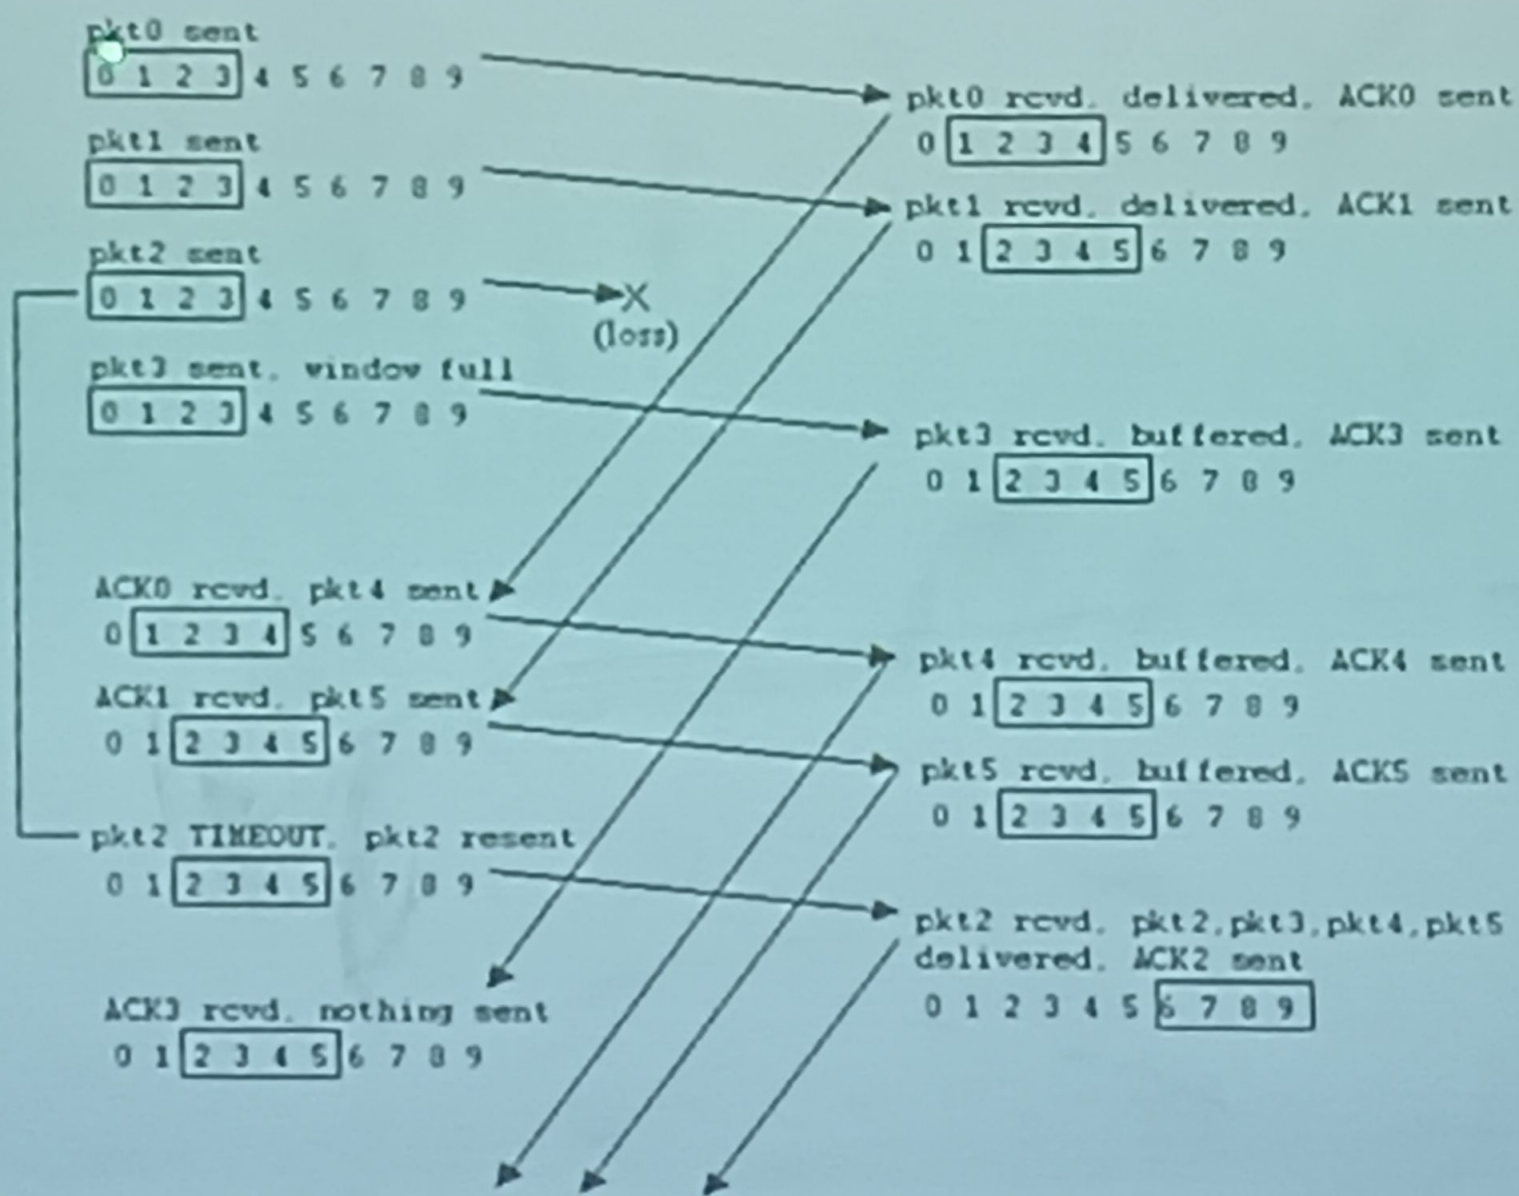
Send window for Selective Repeat ARQ



Receive window for Selective Repeat ARQ



Selective repeat in action



Selective repeat: dilemma

Example:

- seq #'s: 0, 1, 2, 3
- window size=3

- receiver sees no difference in two scenarios!
- incorrectly passes duplicate data as new in (a)

Q: what relationship between seq # size and window size?

